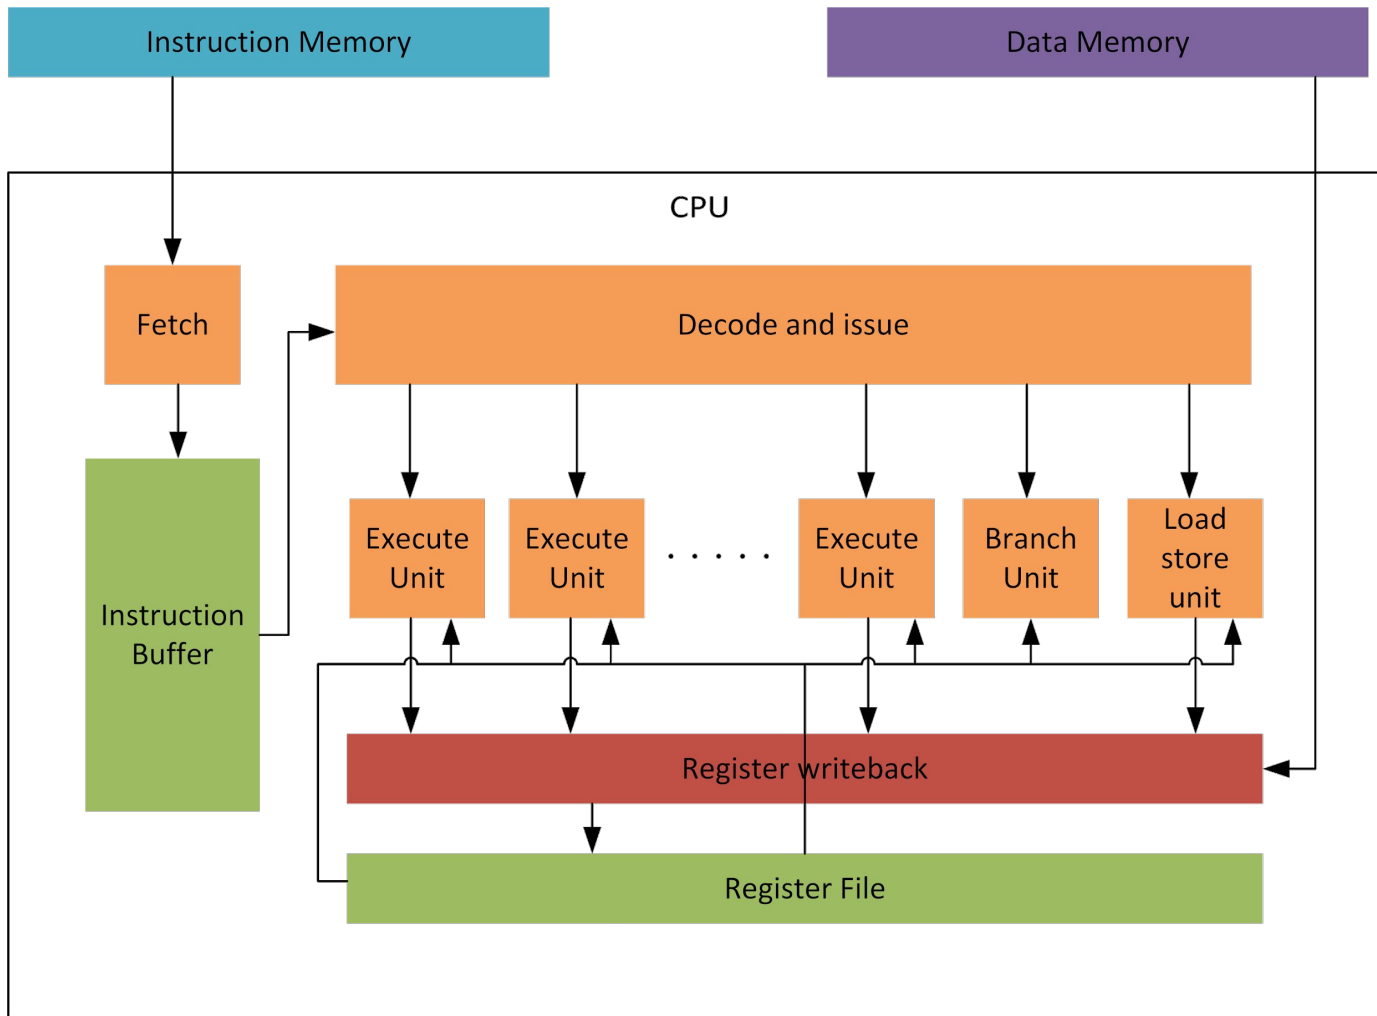


Advanced Computer Architecture

Alex Harman



My Architecture

- N execute units
- Out of order
- 1 load/store unit
- 1 branch unit
- 5-deep pipeline



Simulator options

- -a Assemble program
- -e Number of execution units to use
- -b Size of the instruction buffer
- -f File to use
- -d Debug mode, writing the state every iteration
- -i Interactive mode, steps through cycle by cycle and displays the state
- -h Help, including a full list of options

There are several other options used to tweak the debug output.

Example Program:

Assembly:

```
MEM      0 0
MEM      1 0
MEM      2 0xDEAD
MEM      3 0xDEADBEEF

LDI      R0 2
MVI      R1 0xBEEF
MVI      R2 3
LD        R2 R2

LSI      R0 R0 1      "MyLabel"
ADD      R3 R0 R1
BRNI     R2 R3 >MyLabel

STI      R2 4
TERMINATE
```

Machine Code:

```
00000000000000000000000000000000
00000000000000000000000000000000
000000000000000001101111010101101
110111101010110110111111011101111
---
00100011000000000000000000000010
00100111000100001011111011101111
00100111001000000000000000000011
00100010001000100000000000000000
00001101000000000000000000000001
00000010001100000001000000000000
01000111001000110000000000000100
00101001001000000000000000000100
11111111111111111111111111111111
```

Branching:

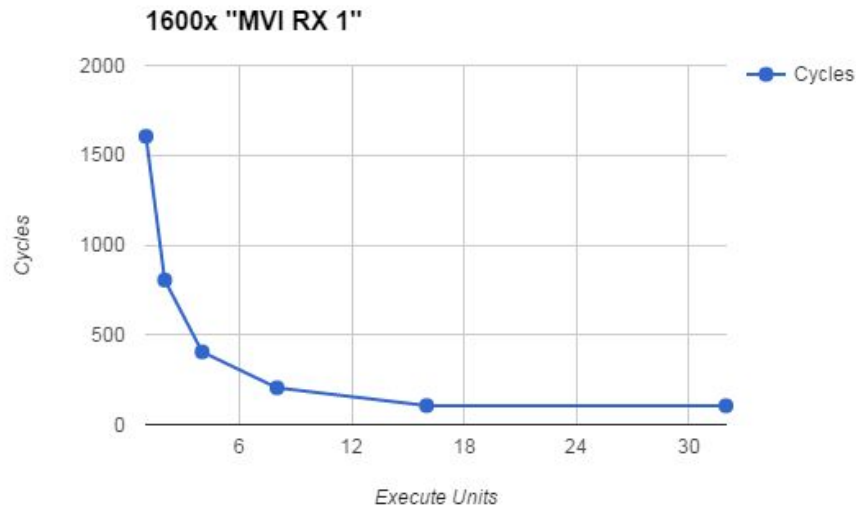
- Simulator has no branch prediction, so branching is expensive
 - To demonstrate this, I use a simple fibonacci sequence calculator up to Fib(47).
 - Best case: 236 cycles with looping, vs 98 unrolled.
 - 1.0 instructions per cycle with looping vs 2.0 unrolled.
-
- The branching version benefitted from superscalar execution, whereas unrolled version did not.
 - With only 1 EU, it took 373 cycles compared to 236 with 3 or more.

Experiment 1 - Performance is capped at 16 EUs

- Because there are only 16 registers, the largest number of independent instructions is 16
- Increasing beyond this amount cannot improve performance, as there is no further parallelism to exploit
- Using, “MVI RX 1” for each register, dependencies are unavoidable after 16 instructions. After using all registers, repeat.
- 1600 instructions (+1 for terminate)

Experiment 1 - Result:

- There is a 1 cycle performance increase between 16 and 32 execute units.
 - This is because the 32 EU version can handle the terminate instruction at the same time as an MVI
- Beyond this, there is no benefit to using more than 16 execute units



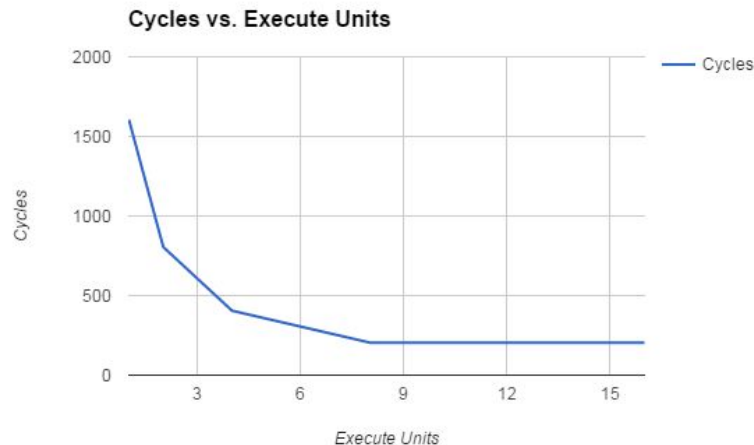
Experiment 2 - Lack of bypassing effective cap

- Because of register writeback, ALU operations take an extra cycle past executing to complete.
- This can effectively cap the number of useful execute units at 8, due to read-after-write dependencies.
- Using only instructions such as “ADDI RX RX 1” we can force bubbles to appear in the pipeline to demonstrate this.

Experiment 2 - Result

As can be seen, the bubbles in the pipeline effectively half the throughput for these programs.

With a cap of 16 independent instructions, but half the pipeline being empty, we get the same performance using 16 EUs as we do with 8.



MEM ACCESS:	0	NOP				MEM ACCESS:	0	ADDI	R0	R0	0x1
	0	NOP					0	ADDI	R1	R1	0x1
	0	NOP					0	ADDI	R2	R2	0x1
	0	NOP					0	ADDI	R3	R3	0x1
EXECUTE:	0	ADDI	R0	R0	0x1	EXECUTE:	0	NOP			
	0	ADDI	R1	R1	0x1		0	NOP			
	0	ADDI	R2	R2	0x1		0	NOP			
	0	ADDI	R3	R3	0x1		0	NOP			
WRITE BACK:	0	NOP				WRITE BACK:	0	ADDI	R0	R0	0x1
	0	NOP					0	ADDI	R1	R1	0x1
	0	NOP					0	ADDI	R2	R2	0x1
	0	NOP					0	ADDI	R3	R3	0x1

Experiment 3: Comparing in-order and out-of-order

- All of these programs do the same thing, they just order their instructions differently.
 - Along the x axis, we increase the proximity of dependent instructions.
- You can see how out-of-order execution allows the processor to scale much better than in-order.
- The in-order processor shows next to no scaling past 2 execute units.

